

# DNF

October 24, 2023

```
[ ]: import pandas as pd
import statsmodels
import numpy as np
import matplotlib.pyplot as plt

status = pd.read_csv('status.csv',encoding = "utf-8")
drivers = pd.read_csv('drivers.csv',encoding = "utf-8")
constructors = pd.read_csv('constructors.csv',encoding = "utf-8")
circuits = pd.read_csv('circuits.csv',encoding = "utf-8")
races = pd.read_csv('races.csv',encoding = "utf-8")
results = pd.read_csv('results.csv', encoding="utf-8")
constructor_standings = pd.read_csv('constructor_standings.csv',encoding = "utf-8")
print('This is the Status data')
status.head()
```

This is the Status data

```
[ ]:      statusId      status
0         1      Finished
1         2  Disqualified
2         3      Accident
3         4      Collision
4         5          Engine
```

```
[ ]: drivers.head()
```

```
[ ]:      driverId  driverRef number code  forename      surname      dob \
0         1      hamilton   44  HAM      Lewis      Hamilton  1985-01-07
1         2      heidfeld   \N  HEI      Nick      Heidfeld  1977-05-10
2         3      rosberg    6   ROS      Nico      Rosberg   1985-06-27
3         4      alonso     14  ALO  Fernando      Alonso   1981-07-29
4         5  kovalainen     \N  KOV      Heikki   Kovalainen  1981-10-19

      nationality      url
0      British      http://en.wikipedia.org/wiki/Lewis_Hamilton
1      German      http://en.wikipedia.org/wiki/Nick_Heidfeld
2      German      http://en.wikipedia.org/wiki/Nico_Rosberg
```

```

3     Spanish    http://en.wikipedia.org/wiki/Fernando_Alonso
4     Finnish    http://en.wikipedia.org/wiki/Heikki_Kovalainen

```

```

[ ]: drivers['driver_name'] = drivers['forename'] + ' ' + drivers['surname']
col = ['driverId', 'driver_name', 'nationality']
my_drivers = drivers[col]
my_drivers.rename(columns = {'nationality': 'driver_nationality' }, inplace =
↳ True)
my_drivers.head()

```

```

/var/folders/zj/rx7n4ybx1ml9886cdbwv429m0000gn/T/ipykernel_84467/3935489618.py:4
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

my_drivers.rename(columns = {'nationality': 'driver_nationality' }, inplace =
True)

```

```

[ ]:
   driverId  driver_name driver_nationality
0         1   Lewis Hamilton           British
1         2   Nick Heidfeld           German
2         3   Nico Rosberg           German
3         4   Fernando Alonso          Spanish
4         5   Heikki Kovalainen          Finnish

```

```

[ ]: constructors.head()

```

```

[ ]:
   constructorId  constructorRef      name nationality \
0              1      mclaren    McLaren    British
1              2    bmw_sauber  BMW Sauber    German
2              3    williams   Williams    British
3              4    renault    Renault    French
4              5    toro_rosso  Toro Rosso    Italian

```

```

                                url
0      http://en.wikipedia.org/wiki/McLaren
1      http://en.wikipedia.org/wiki/BMW_Sauber
2  http://en.wikipedia.org/wiki/Williams_Grand_Pr...
3  http://en.wikipedia.org/wiki/Renault_in_Formul...
4  http://en.wikipedia.org/wiki/Scuderia_Toro_Rosso

```

```

[ ]: cols = ['constructorId', 'name', 'nationality']
my_constructors = constructors[cols]
my_constructors.rename(columns = {'nationality': 'team_nationality' }, inplace =
↳ True)
my_constructors.rename(columns = {'name': 'constructor' }, inplace = True)
my_constructors.head()

```

```
/var/folders/zj/rx7n4ybx1ml9886cdbwv429m0000gn/T/ipykernel_84467/1242970380.py:3
: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
my_constructors.rename(columns = {'nationality': 'team_nationality' }, inplace
= True)
```

```
/var/folders/zj/rx7n4ybx1ml9886cdbwv429m0000gn/T/ipykernel_84467/1242970380.py:4
: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
my_constructors.rename(columns = {'name': 'constructor' }, inplace = True)
```

```
[ ]: constructorId constructor team_nationality
0          1      McLaren      British
1          2    BMW Sauber      German
2          3    Williams      British
3          4      Renault      French
4          5    Toro Rosso      Italian
```

```
[ ]: circuits.head()
```

```
[ ]: circuitId  circuitRef      name  location \
0          1  albert_park  Albert Park Grand Prix Circuit  Melbourne
1          2      sepang   Sepang International Circuit  Kuala Lumpur
2          3    bahrain   Bahrain International Circuit      Sakhir
3          4  catalunya  Circuit de Barcelona-Catalunya  Montmeló
4          5    istanbul              Istanbul Park      Istanbul
```

```
country      lat      lng  alt \
0  Australia -37.84970  144.96800  10
1  Malaysia  2.76083  101.73800  18
2   Bahrain  26.03250  50.51060   7
3    Spain  41.57000   2.26111  109
4    Turkey  40.95170  29.40500  130
```

```
url
0  http://en.wikipedia.org/wiki/Melbourne_Grand_P...
1  http://en.wikipedia.org/wiki/Sepang_Internatio...
2  http://en.wikipedia.org/wiki/Bahrain_Internati...
3  http://en.wikipedia.org/wiki/Circuit_de_Barcel...
4      http://en.wikipedia.org/wiki/Istanbul_Park
```

```
[ ]: colss = ['circuitId', 'circuitRef', 'name', 'location', 'country']
my_circuits = circuits[colss]
my_circuits.rename(columns = {'name': 'circuit' }, inplace = True)
my_circuits.head()
```

/var/folders/zj/rx7n4ybx1ml9886cdbwv429m0000gn/T/ipykernel\_84467/823439227.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
my_circuits.rename(columns = {'name': 'circuit' }, inplace = True)
```

```
[ ]: circuitId  circuitRef      circuit      location \
0           1  albert_park  Albert Park Grand Prix Circuit      Melbourne
1           2      sepang    Sepang International Circuit      Kuala Lumpur
2           3    bahrain    Bahrain International Circuit      Sakhir
3           4  catalunya  Circuit de Barcelona-Catalunya      Montmeló
4           5    istanbul      Istanbul Park      Istanbul

      country
0  Australia
1  Malaysia
2   Bahrain
3     Spain
4     Turkey
```

```
[ ]: constructor_standings.head()
```

```
[ ]: constructorStandingsId  raceId  constructorId  points  position \
0                        1      18              1    14.0         1
1                        2      18              2     8.0         3
2                        3      18              3     9.0         2
3                        4      18              4     5.0         4
4                        5      18              5     2.0         5

      positionText  wins
0                1     1
1                3     0
2                2     0
3                4     0
4                5     0
```

```
[ ]: colss4 = ['constructorStandingsId', 'raceId', 'constructorId', 'points']
my_constructor_standings = constructor_standings[colss4]
my_constructor_standings.head()
```

```
[ ]: constructorStandingsId  raceId  constructorId  points
0          1          18          1      14.0
1          2          18          2       8.0
2          3          18          3       9.0
3          4          18          4       5.0
4          5          18          5       2.0
```

```
[ ]: races.head()
```

```
[ ]: raceId  year  round  circuitId          name          date \
0         1  2009     1           1  Australian Grand Prix  2009-03-29
1         2  2009     2           2   Malaysian Grand Prix  2009-04-05
2         3  2009     3          17   Chinese Grand Prix  2009-04-19
3         4  2009     4           3   Bahrain Grand Prix  2009-04-26
4         5  2009     5           4   Spanish Grand Prix  2009-05-10

      time                                     url fp1_date \
0  06:00:00  http://en.wikipedia.org/wiki/2009_Australian_G...  \N
1  09:00:00  http://en.wikipedia.org/wiki/2009_Malaysian_Gr...  \N
2  07:00:00  http://en.wikipedia.org/wiki/2009_Chinese_Gran...  \N
3  12:00:00  http://en.wikipedia.org/wiki/2009_Bahrain_Gran...  \N
4  12:00:00  http://en.wikipedia.org/wiki/2009_Spanish_Gran...  \N

fp1_time fp2_date fp2_time fp3_date fp3_time quali_date quali_time \
0      \N      \N      \N      \N      \N      \N      \N
1      \N      \N      \N      \N      \N      \N      \N
2      \N      \N      \N      \N      \N      \N      \N
3      \N      \N      \N      \N      \N      \N      \N
4      \N      \N      \N      \N      \N      \N      \N

sprint_date sprint_time
0      \N      \N
1      \N      \N
2      \N      \N
3      \N      \N
4      \N      \N
```

```
[ ]: colsss = ['raceId', 'year', 'round', 'circuitId', 'name']
my_races = races[colsss]
my_races.rename(columns = {'name': 'prix' }, inplace = True)
my_races.head()
```

```
/var/folders/zj/rx7n4ybx1ml9886cdbwv429m0000gn/T/ipykernel_84467/1917903944.py:3
: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
my_races.rename(columns = {'name': 'prix' }, inplace = True)
```

```
[ ]:   raceId  year  round  circuitId      prix
0      1  2009    1         1  Australian Grand Prix
1      2  2009    2         2   Malaysian Grand Prix
2      3  2009    3        17   Chinese Grand Prix
3      4  2009    4         3   Bahrain Grand Prix
4      5  2009    5         4   Spanish Grand Prix
```

```
[ ]: results.head()
```

```
[ ]:   resultId  raceId  driverId  constructorId  number  grid position \
0          1      18         1             1      22      1         1
1          2      18         2             2       3       5         2
2          3      18         3             3       7       7         3
3          4      18         4             4       5      11         4
4          5      18         5             1      23       3         5

   positionText  positionOrder  points  laps      time milliseconds \
0             1              1    10.0   58  1:34:50.616      5690616
1             2              2     8.0   58      +5.478      5696094
2             3              3     6.0   58      +8.163      5698779
3             4              4     5.0   58     +17.181      5707797
4             5              5     4.0   58     +18.014      5708630

   fastestLap  rank  fastestLapTime  fastestLapSpeed  statusId
0          39    2      1:27.452         218.300          1
1          41    3      1:27.739         217.586          1
2          41    5      1:28.090         216.719          1
3          58    7      1:28.603         215.464          1
4          43    1      1:27.418         218.385          1
```

```
[ ]: col2 = ['resultId', 'raceId', 'driverId', 'constructorId', 'points', 'laps', 'statusId']
my_results = results[col2]
my_results.head()
```

```
[ ]:   resultId  raceId  driverId  constructorId  points  laps  statusId
0          1      18         1             1    10.0   58          1
1          2      18         2             2     8.0   58          1
2          3      18         3             3     6.0   58          1
3          4      18         4             4     5.0   58          1
4          5      18         5             1     4.0   58          1
```

```
[ ]: # make one dataframe
my_data = pd.merge(my_results, my_races, on = 'raceId', how = 'left')
my_data = pd.merge(my_data, my_circuits, on = 'circuitId', how = 'left')
my_data = pd.merge(my_data, my_drivers, on = 'driverId', how = 'left')
```

```

my_data = pd.merge(my_data, status, on = 'statusId', how = 'left')
my_data = pd.merge(my_data, my_constructors, on = 'constructorId', how = 'left')
my_data = pd.merge(my_data, my_constructor_standings, on = ['constructorId', 'raceId'], how = 'left')
my_data.head()

```

```

[ ]:
  resultId  raceId  driverId  constructorId  points_x  laps  statusId  year \
0         1      18         1              1      10.0   58         1  2008
1         2      18         2              2       8.0   58         1  2008
2         3      18         3              3       6.0   58         1  2008
3         4      18         4              4       5.0   58         1  2008
4         5      18         5              1       4.0   58         1  2008

```

```

      round  circuitId  ...      circuit  location \
0         1          1  ...  Albert Park Grand Prix Circuit  Melbourne
1         1          1  ...  Albert Park Grand Prix Circuit  Melbourne
2         1          1  ...  Albert Park Grand Prix Circuit  Melbourne
3         1          1  ...  Albert Park Grand Prix Circuit  Melbourne
4         1          1  ...  Albert Park Grand Prix Circuit  Melbourne

```

```

      country      driver_name  driver_nationality  status  constructor \
0  Australia  Lewis Hamilton      British  Finished    McLaren
1  Australia  Nick Heidfeld      German  Finished    BMW Sauber
2  Australia  Nico Rosberg      German  Finished    Williams
3  Australia  Fernando Alonso    Spanish  Finished    Renault
4  Australia  Heikki Kovalainen  Finnish  Finished    McLaren

```

```

      team_nationality  constructorStandingsId  points_y
0         British              1.0      14.0
1         German              2.0       8.0
2         British              3.0       9.0
3         French              4.0       5.0
4         British              1.0      14.0

```

[5 rows x 22 columns]

```

[ ]: # Dataframe of all drivers who finished a Grand Prix
finished_race = my_data.loc[my_data['status'].isin(['Finished', '+1 Lap', '+2 Laps', '+3 Laps', '+4 Laps', '+5 Laps', '+6 Laps', '+7 Laps', '+8 Laps', '+9 Laps', '+10 Laps', '+11 Laps', '+12 Laps', '+13 Laps', '+14 Laps'])] # included the overlapped status

# Percentage of drivers finishing a race in all of F1 history
finish = round((finished_race.shape[0] / my_data.shape[0]) * 100, 3)

```

```
print(f"In all of F1 history, {finish}% drivers have finished a Grand Prix that_\nthey started")
```

In all of F1 history, 55.691% drivers have finished a Grand Prix that they started

```
[ ]: my_data.describe()
```

```
[ ]:
```

	resultId	raceId	driverId	constructorId	points_x \
count	26180.000000	26180.000000	26180.000000	26180.000000	26180.000000
mean	13091.388426	538.896982	268.148739	49.237051	1.918795
std	7558.893818	304.537441	274.174971	60.420088	4.240048
min	1.000000	1.000000	1.000000	1.000000	0.000000
25%	6545.750000	295.000000	57.000000	6.000000	0.000000
50%	13090.500000	521.000000	164.000000	25.000000	0.000000
75%	19635.250000	794.000000	373.000000	59.000000	2.000000
max	26185.000000	1115.000000	859.000000	214.000000	50.000000

	laps	statusId	year	round	circuitId \
count	26180.000000	26180.000000	26180.000000	26180.000000	26180.000000
mean	46.101795	17.459358	1990.677082	8.393965	23.579908
std	29.691532	26.161846	19.573972	4.957933	18.782538
min	0.000000	1.000000	1950.000000	1.000000	1.000000
25%	22.000000	1.000000	1976.000000	4.000000	9.000000
50%	53.000000	10.000000	1991.000000	8.000000	18.000000
75%	66.000000	14.000000	2008.000000	12.000000	34.000000
max	200.000000	141.000000	2023.000000	22.000000	79.000000

	constructorStandingsId	points_y
count	24313.000000	24313.000000
mean	16719.382224	38.125756
std	8916.962800	80.463991
min	1.000000	0.000000
25%	8685.000000	1.000000
50%	20138.000000	10.000000
75%	25045.000000	36.000000
max	28632.000000	765.000000

```
[ ]: # List of statusId values to exclude - those who finished
exclude_ids = [1, 11, 12, 13, 14, 15, 16, 17, 18, 19, 45, 50, 128, 53, 55, 58,\n88, 111, 112,\n113, 114, 115, 116, 117, 118, 119, 120, 122, 123, 124, 125, 127,\n133, 134]

# Filter data to exclude specific statusId values and include all the rest
crashes = my_data.loc[~my_data['statusId'].isin(exclude_ids)]
```



```
# Describe the resulting data
crashes.describe()
#11524
```

```
[ ]:      resultId      raceId      driverId  constructorId      points_x \
count  11524.000000  11524.000000  11524.000000   11524.000000  11524.000000
mean    11865.126605    494.742798    238.241843     51.008330    0.013016
std     6353.839707    249.597802    219.073648     55.302585    0.187255
min       7.000000     1.000000     1.000000     1.000000    0.000000
25%     6700.750000    308.000000     88.000000    15.000000    0.000000
50%    11589.500000    477.000000    170.000000    32.000000    0.000000
75%    16871.000000    681.000000    327.250000    60.000000    0.000000
max    26185.000000   1115.000000    858.000000   214.000000    6.000000

      laps      statusId      year      round      circuitId \
count  11524.000000  11524.000000  11524.000000  11524.000000  11524.000000
mean     24.274731    30.516314  1985.009198     7.826189    23.992103
std     24.745933    33.581783    17.070282     4.650605    17.857834
min       0.000000     2.000000  1950.000000     1.000000     1.000000
25%       2.000000     5.000000  1974.000000     4.000000    10.000000
50%      19.000000    10.000000  1986.000000     7.000000    19.000000
75%      40.000000    54.000000  1996.000000    11.000000    36.000000
max     196.000000   141.000000  2023.000000    22.000000    79.000000

      constructorStandingsId      points_y
count          10479.000000  10479.000000
mean           15096.260807    18.181649
std            8182.478043    44.405906
min              2.000000     0.000000
25%            8487.000000     0.000000
50%           11282.000000     4.000000
75%           22865.500000    17.000000
max           28629.000000   701.000000
```

```
[ ]: # Remove duplicates based on 'raceId' and 'driverId'
crashes_unique = crashes.drop_duplicates(subset=['raceId', 'driverId'])

crashes_unique.describe()
```

```
[ ]:      resultId      raceId      driverId  constructorId      points_x \
count  11498.000000  11498.000000  11498.000000   11498.000000  11498.000000
mean    11847.029396    494.093147    237.580275     50.890503    0.013046
std     6349.071231    249.491276    218.839733     55.278920    0.187466
min       7.000000     1.000000     1.000000     1.000000    0.000000
25%     6694.250000    308.000000     88.000000    15.000000    0.000000
50%    11563.500000    476.000000    170.000000    32.000000    0.000000
75%    16841.750000    679.000000    326.000000    59.000000    0.000000
```

max	26185.000000	1115.000000	858.000000	214.000000	6.000000
-----	--------------	-------------	------------	------------	----------

	laps	statusId	year	round	circuitId \
count	11498.000000	11498.000000	11498.000000	11498.000000	11498.000000
mean	24.143416	30.504175	1985.074709	7.835624	23.988607
std	24.425842	33.576483	17.031860	4.649396	17.858035
min	0.000000	2.000000	1950.000000	1.000000	1.000000
25%	2.000000	5.000000	1974.000000	4.000000	10.000000
50%	19.000000	10.000000	1986.000000	7.000000	19.000000
75%	40.000000	54.000000	1996.000000	11.000000	36.000000
max	196.000000	141.000000	2023.000000	22.000000	79.000000

	constructorStandingsId	points_y
count	10473.000000	10473.000000
mean	15093.612336	18.184713
std	8183.935630	44.417275
min	2.000000	0.000000
25%	8485.000000	0.000000
50%	11280.000000	4.000000
75%	22866.000000	17.000000
max	28629.000000	701.000000

```
[ ]: crashes.head()
```

```
[ ]:
```

	resultId	raceId	driverId	constructorId	points_x	laps	statusId	year	\
6	7	18	7	5	2.0	55	5	2008	
7	8	18	8	6	1.0	53	5	2008	
8	9	18	9	2	0.0	47	4	2008	
9	10	18	10	7	0.0	43	3	2008	
10	11	18	11	8	0.0	32	7	2008	

	round	circuitId	...	circuit	location	\
6	1	1	...	Albert Park Grand Prix Circuit	Melbourne	
7	1	1	...	Albert Park Grand Prix Circuit	Melbourne	
8	1	1	...	Albert Park Grand Prix Circuit	Melbourne	
9	1	1	...	Albert Park Grand Prix Circuit	Melbourne	
10	1	1	...	Albert Park Grand Prix Circuit	Melbourne	

	country	driver_name	driver_nationality	status	\
6	Australia	Sébastien Bourdais	French	Engine	
7	Australia	Kimi Räikkönen	Finnish	Engine	
8	Australia	Robert Kubica	Polish	Collision	
9	Australia	Timo Glock	German	Accident	
10	Australia	Takuma Sato	Japanese	Transmission	

	constructor	team_nationality	constructorStandingsId	points_y
6	Toro Rosso	Italian	5.0	2.0

7	Ferrari	Italian	6.0	1.0
8	BMW Sauber	German	2.0	8.0
9	Toyota	Japanese	NaN	NaN
10	Super Aguri	Japanese	NaN	NaN


[5 rows x 22 columns]


```
[ ]: def plot_crashes_by_columns(crashes, column_names):
    """Plot the number of crashes based on the selected columns from the
    ↪crashes df."""


    for column_name in column_names:
        crashes_by_column = crashes.groupby(column_name).size()

        plt.figure(figsize=(12, 10))
        plt.bar(crashes_by_column.index, crashes_by_column.values)
        plt.title(f'Number of Crashes by {column_name}')
        plt.ylabel('Number of Crashes')
        plt.xlabel(column_name)
        plt.xticks(crashes_by_column.index, rotation=90)
        plt.tight_layout()
        plt.show()

# Example usage:
columns_to_plot = ['year', 'circuit', 'prix']
plot_crashes_by_columns(crashes, columns_to_plot)
```








```
[ ]: DNF_year = crashes.groupby('year').size()

plt.figure(figsize=(12, 10))
plt.plot(DNF_year.index, DNF_year.values)
plt.title(f'Number of Race Retirements by Year')
plt.ylabel('Number of Race Retirements')
plt.xlabel('Year')
plt.xticks(DNF_year.index, rotation=90)
plt.tight_layout()
plt.show()
```



```
[ ]: import plotly.express as px
import plotly.graph_objects as go

def plot_dnf(year):
    # Filter crashes for the given year
    crashes_in_specific_year = crashes[crashes['year'] == year]

    # For Drivers
    crashes_driver = crashes_in_specific_year.groupby('driver_name').size().
    ↪reset_index(name='counts')
    scaled_size_driver = crashes_driver['counts']**2
    fig_driver = create_bubble_chart(crashes_driver, 'driver_name',
    ↪scaled_size_driver, 'Driver Name', f'DNFs by Driver in {year}')
    fig_driver.show()

    # For Constructors
    crashes_team = crashes_in_specific_year.groupby('constructor').size().
    ↪reset_index(name='counts')
```

```

scaled_size_team = crashes_team['counts']**2
fig_team = create_bubble_chart(crashes_team, 'constructor',
↪scaled_size_team, 'Constructor', f'DNFs by Constructor in {year}')
fig_team.show()

def create_bubble_chart(df, x_col, scaled_size, xaxis_title, chart_title):
    colors = px.colors.qualitative.Plotly

    fig = go.Figure(data=[go.Scatter(
        x=df[x_col], y=df['counts'],
        mode='markers',
        marker=dict(
            color=colors * (len(df) // len(colors)) + colors[:len(df) %
↪len(colors)],
            size=scaled_size / max(scaled_size) * 100, # Normalize and then
↪scale to desired range
            sizemin=6 # Minimum marker size
        )
    )])

    fig.update_layout(
        height=800,
        yaxis=dict(title='Number of DNFs'),
        xaxis=dict(title=xaxis_title),
        title=chart_title
    )

    return fig

```

```

[ ]: # Use the function
plot_dnf(2021)

```

```

[ ]: races.head()
races_points = races[['raceId', 'year']]
races_points.head()

```

```

[ ]:
  raceId  year
0      1  2009
1      2  2009
2      3  2009
3      4  2009
4      5  2009

```

```

[ ]: constructor_standings.head()
constructor_points = constructor_standings[['constructorStandingsId', 'raceId',
↪'constructorId', 'points']]
constructor_points.head()

```



```
[ ]: constructorStandingsId  raceId  constructorId  points
0                1         18                1    14.0
1                2         18                2     8.0
2                3         18                3     9.0
3                4         18                4     5.0
4                5         18                5     2.0
```

```
[ ]: constructors.head()
cons = constructors[['constructorId', 'name']]
cons.head()
```

```
[ ]: constructorId      name
0          1    McLaren
1          2  BMW Sauber
2          3   Williams
3          4   Renault
4          5  Toro Rosso
```

```
[ ]: final_df = pd.merge(constructor_points, cons, on = 'constructorId', how = 'left')
final_df = pd.merge(final_df, races_points, on = 'raceId', how = 'left')
final_df.head()
```

```
[ ]: constructorStandingsId  raceId  constructorId  points      name  year
0                1         18                1    14.0    McLaren  2008
1                2         18                2     8.0  BMW Sauber  2008
2                3         18                3     9.0   Williams  2008
3                4         18                4     5.0   Renault  2008
4                5         18                5     2.0  Toro Rosso  2008
```

```
[ ]: final_df['total_points'] = final_df.groupby(['raceId', 'constructorId'])['points'].transform('sum')

# Since we have a sum, drop duplicates
last_df = final_df[['constructorId', 'points', 'name', 'year']]
last_df = last_df.drop_duplicates(subset=['constructorId', 'year'])
```

```
[ ]: constructorId  points      name  year
12316          210     1.0  Haas F1 Team  2020
12317           9    78.0    Red Bull  2020
12318           3     0.0    Williams  2020
12319          51     2.0   Alfa Romeo  2020
12320           4    32.0    Renault  2020
12321         213    13.0  AlphaTauri  2020
12322         211    42.0  Racing Point  2020
12323           1    51.0    McLaren  2020
12324           6    43.0    Ferrari  2020
```

12325	131	146.0	Mercedes	2020
-------	-----	-------	----------	------

```
[ ]: filtered_df = last_df[last_df['year'] == 2020]  
filtered_df
```